

✓ Baseline

```

import torch
import torchvision
import numpy as np
import matplotlib.pyplot as plt
import PIL

from torchsummary import summary

# From local helper files
from helper_evaluation import set_all_seeds, set_deterministic, compute_confusion_matrix
from helper_train import train_model
from helper_plotting import plot_training_loss, plot_accuracy, show_examples, plot_confusion_matrix
from helper_dataset import get_dataloaders_cifar10, UnNormalize

RANDOM_SEED = 123
BATCH_SIZE = 256
NUM_EPOCHS = 40
DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

train_transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((16, 16)),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

test_transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((16, 16)),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

train_loader, valid_loader, test_loader = get_dataloaders_cifar10(
    batch_size=BATCH_SIZE,
    validation_fraction=0.1,
    train_transforms=train_transforms,
    test_transforms=test_transforms,
    num_workers=2)

# Checking the dataset
for images, labels in train_loader:
    print('Image batch dimensions:', images.shape)
    print('Image label dimensions:', labels.shape)
    print('Class labels of 10 examples:', labels[:10])
    break

📄 Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to data/cifar-10-python.tar.gz
100% 170498071/170498071 [00:02<00:00, 72251698.13it/s]
Extracting data/cifar-10-python.tar.gz to data
Image batch dimensions: torch.Size([256, 3, 16, 16])
Image label dimensions: torch.Size([256])

class CNN1(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.features = torch.nn.Sequential(
            # Conv 1
            torch.nn.Conv2d(3, 16, kernel_size=3, padding="same"), # output 16 - 3 + 1 => 16
            # , stride=4, padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2), # 16 / 2 => output 8

            # Conv 2
            torch.nn.Conv2d(16, 32, kernel_size=3, padding="same"), # output 7 - 2 + 1 => 8
            # , padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2) #output 8 / 2 => output 4

```

```

)

self.classifier = torch.nn.Sequential(
    torch.nn.Linear(32*4*4, 100),
    torch.nn.ReLU(inplace=True),
    torch.nn.Linear(100, num_classes),
)

def forward(self, x):
    x = self.features(x)
    x = torch.flatten(x, 1)
    # print(x.size())
    logits = self.classifier(x)
    return logits

```

```
model1 = CNN1(num_classes=10)
```

```
model1 = model1.to(DEVICE)
```

```
print(summary(model1, (3, 16, 16)))
```

```

-----
Layer (type)          Output Shape          Param #
-----
Conv2d-1              [-1, 16, 16, 16]      448
ReLU-2                [-1, 16, 16, 16]      0
MaxPool2d-3           [-1, 16, 8, 8]        0
Conv2d-4              [-1, 32, 8, 8]        4,640
ReLU-5                [-1, 32, 8, 8]        0
MaxPool2d-6           [-1, 32, 4, 4]        0
Linear-7              [-1, 100]             51,300
ReLU-8               [-1, 100]             0
Linear-9              [-1, 10]              1,010
-----
Total params: 57,398
Trainable params: 57,398
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.11
Params size (MB): 0.22
Estimated Total Size (MB): 0.33
-----
None

```

```

optimizer1 = torch.optim.SGD(model1.parameters(), lr=0.1)
scheduler1 = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer,
                                                         factor=0.1,
                                                         mode='max',
                                                         verbose=True)

```

```

minibatch_loss_list1, train_acc_list1, valid_acc_list1 = train_model(
    model=model1,
    num_epochs=NUM_EPOCHS,
    train_loader=train_loader,
    valid_loader=valid_loader,
    test_loader=test_loader,
    optimizer=optimizer1,
    device=DEVICE,
    scheduler=None,
    scheduler_on='valid_acc',
    logging_interval=100)

```

```

plot_training_loss(minibatch_loss_list=minibatch_loss_list1,
                  num_epochs=NUM_EPOCHS,
                  iter_per_epoch=len(train_loader),
                  results_dir=None,
                  averaging_iterations=20)

```


```
plt.show()
```

```

plot_accuracy(train_acc_list=train_acc_list1,
              valid_acc_list=valid_acc_list1,
              results_dir=None)

```

```
# plt.ylim([80, 100])
plt.show()
```



```
class CNN2(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.features = torch.nn.Sequential(
            # Conv 1
            torch.nn.Conv2d(3, 16, kernel_size=3, padding="same"), # output 16 - 3 + 1 => 16
            # , stride=4, padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2), # 16 / 2 => output 8

            # Conv 2
            torch.nn.Conv2d(16, 32, kernel_size=2, padding="same"), # output 7 - 2 + 1 => 8
            # , padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2), #output 8 / 2 => output 4

            # Conv 3
            torch.nn.Conv2d(32, 64, kernel_size=2, padding="same"), # output 7 - 2 + 1 => 4
            # , padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2) #output 4 / 2 => output 2

        )

        self.classifier = torch.nn.Sequential(
            torch.nn.Linear(64*2*2, 100),
            torch.nn.ReLU(inplace=True),
            torch.nn.Linear(100, num_classes),
        )

    def forward(self, x):
        x = self.features(x)
        x = torch.flatten(x, 1)
```

```
# print(x.size())
logits = self.classifier(x)
return logits
```

```
model2 = CNN2(num_classes=10)
```

```
model2 = model2.to(DEVICE)
```

```
print(summary(model2, (3, 16, 16)))
```

```
-----
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 16, 16]	448
ReLU-2	[-1, 16, 16, 16]	0
MaxPool2d-3	[-1, 16, 8, 8]	0
Conv2d-4	[-1, 32, 8, 8]	2,080
ReLU-5	[-1, 32, 8, 8]	0
MaxPool2d-6	[-1, 32, 4, 4]	0
Conv2d-7	[-1, 64, 4, 4]	8,256
ReLU-8	[-1, 64, 4, 4]	0
MaxPool2d-9	[-1, 64, 2, 2]	0
Linear-10	[-1, 100]	25,700
ReLU-11	[-1, 100]	0
Linear-12	[-1, 10]	1,010

```
-----
Total params: 37,494
Trainable params: 37,494
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 0.12
Params size (MB): 0.14
Estimated Total Size (MB): 0.27
-----
None
/usr/local/lib/python3.7/dist-packages/torch/nn/modules/conv.py:454: UserWarning: Using padding='same' with even kernel leng
self.padding, self.dilation, self.groups)
```

```
optimizer2 = torch.optim.SGD(model2.parameters(), lr=0.1)
scheduler2 = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer2,
                                                         factor=0.1,
                                                         mode='max',
                                                         verbose=True)
```


```
minibatch_loss_list2, train_acc_list2, valid_acc_list2 = train_model(
    model=model2,
    num_epochs=NUM_EPOCHS,
    train_loader=train_loader,
    valid_loader=valid_loader,
    test_loader=test_loader,
    optimizer=optimizer2,
    device=DEVICE,
    scheduler=None,
    scheduler_on='valid_acc',
    logging_interval=100)
```

```
plot_training_loss(minibatch_loss_list=minibatch_loss_list2,
                  num_epochs=NUM_EPOCHS,
                  iter_per_epoch=len(train_loader),
                  results_dir=None,
                  averaging_iterations=20)
```

```
plt.show()
```

```
plot_accuracy(train_acc_list=train_acc_list2,
              valid_acc_list=valid_acc_list2,
              results_dir=None)
```

```
# plt.ylim([80, 100])
plt.show()
```



```
import pandas as pd
```

```
results = pd.DataFrame({"Number of Parameters": [57398, 37494], "Accuracy": [80, 70]}, index = ["CNN1", "CNN2"])
```

```
results
```



	Number of Parameters	Accuracy
CNN1	57398	80
CNN2	37494	70

- ✓ The simpler CNN is better for this dataset as the more complex model will have very very less pixels to work with

Start coding or [generate](#) with AI.