

## ✓ Optimiser 1

```
import torch
import torchvision
import numpy as np
import matplotlib.pyplot as plt
import PIL
from torchsummary import summary

# From local helper files
from helper_evaluation import set_all_seeds, set_deterministic, compute_confusion_ma
from helper_train import train_model
from helper_plotting import plot_training_loss, plot_accuracy, show_examples, plot_c
from helper_dataset import get_dataloaders_cifar10, UnNormalize

RANDOM_SEED = 123
BATCH_SIZE = 256
NUM_EPOCHS = 40
DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

```
train_transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((16, 16)),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
```

```
test_transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((16, 16)),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```
train_loader, valid_loader, test_loader = get_dataloaders_cifar10(
    batch_size=BATCH_SIZE,
    validation_fraction=0.1,
    train_transforms=train_transforms,
    test_transforms=test_transforms,
    num_workers=2)
```

# Checking the dataset

```
for images, labels in train_loader:
    print('Image batch dimensions:', images.shape)
    print('Image label dimensions:', labels.shape)
    print('Class labels of 10 examples:', labels[:10])
    break
```

📄 Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to data/cifar-10-python.tar.gz  
100% 170498071/170498071 [00:10<00:00, 17861031.45it/s]

```
Extracting data/cifar-10-python.tar.gz to data
Image batch dimensions: torch.Size([256, 3, 16, 16])
Image label dimensions: torch.Size([256])
```

```
class CNN1RMS(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.features = torch.nn.Sequential(
            # Conv 1
            torch.nn.Conv2d(3, 16, kernel_size=3, padding="same"), # output 16 - 3 +
            # , stride=4, padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2), # 16 / 2 => output 8

            # Conv 2
            torch.nn.Conv2d(16, 32, kernel_size=2, padding="same"), # output 7 - 2 +
            # , padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2) #output 8 / 2 => output 4
        )
```

```

self.classifier = torch.nn.Sequential(
    torch.nn.Linear(32*4*4, 100),
    torch.nn.ReLU(inplace=True),
    torch.nn.Linear(100, num_classes),
)

```

```

def forward(self, x):
    x = self.features(x)
    x = torch.flatten(x, 1)
    # print(x.size())
    logits = self.classifier(x)
    return logits

```

```
model1_rms = CNN1RMS(num_classes=10)
```

```
model1_rms = model1_rms.to(DEVICE)
```

```
print(summary(model1_rms, (3, 16, 16)))
```



Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 16, 16]	448
ReLU-2	[-1, 16, 16, 16]	0
MaxPool2d-3	[-1, 16, 8, 8]	0
Conv2d-4	[-1, 32, 8, 8]	2,080
ReLU-5	[-1, 32, 8, 8]	0
MaxPool2d-6	[-1, 32, 4, 4]	0
Linear-7	[-1, 100]	51,300
ReLU-8	[-1, 100]	0
Linear-9	[-1, 10]	1,010

```

Total params: 54,838
Trainable params: 54,838
Non-trainable params: 0

```

```

Input size (MB): 0.00
Forward/backward pass size (MB): 0.11
Params size (MB): 0.21
Estimated Total Size (MB): 0.32

```

None

```

/usr/local/lib/python3.7/dist-packages/torch/nn/modules/conv.py:454: UserWarning
  self.padding, self.dilation, self.groups)

```

```

optimizer_rms = torch.optim.RMSprop(model1_rms.parameters())
scheduler_rms = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer_rms,
                                                             factor=0.1,

```

```
mode='max',  
verbose=True)
```

```
minibatch_loss_list_rms, train_acc_list_rms, valid_acc_list_rms = train_model(  
    model=model1_rms,  
    num_epochs=NUM_EPOCHS,  
    train_loader=train_loader,  
    valid_loader=valid_loader,  
    test_loader=test_loader,  
    optimizer=optimizer_rms,  
    device=DEVICE,  
    scheduler=None,  
    scheduler_on='valid_acc',  
    logging_interval=100)
```



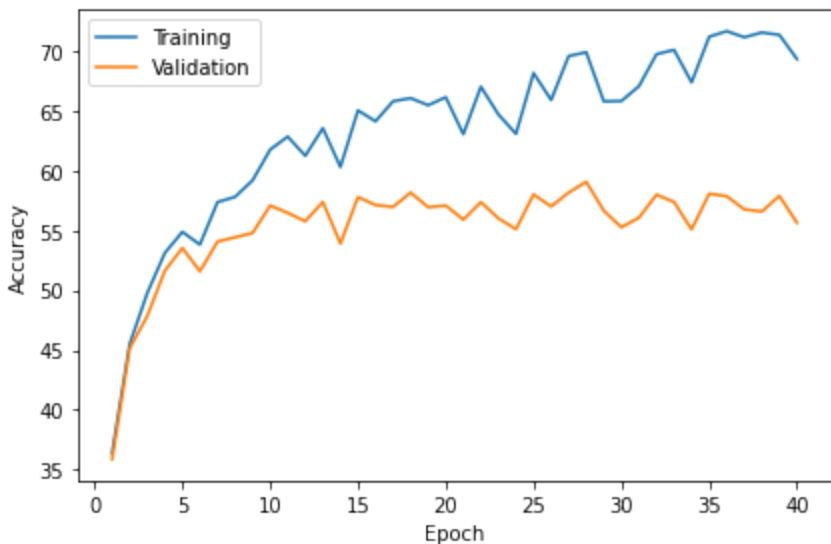
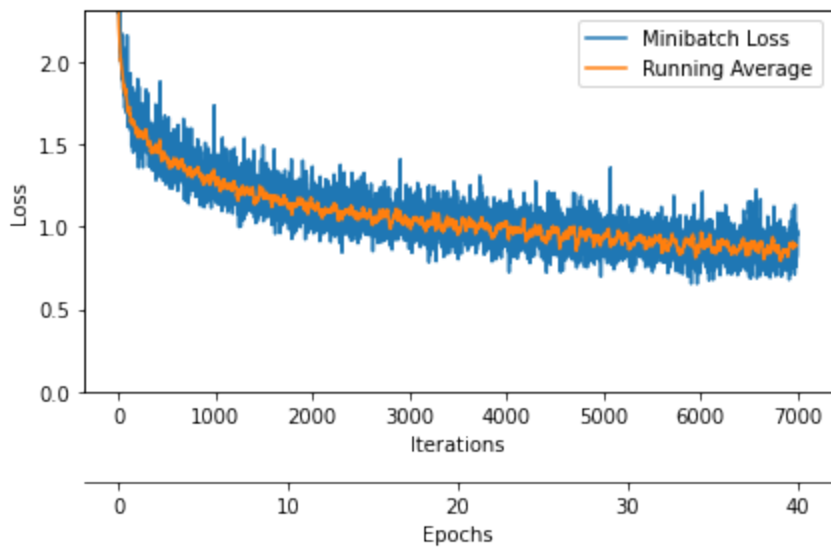
```
Epoch: 037/040 | Batch 0000/0175 | Loss: 0.8526
Epoch: 037/040 | Batch 0100/0175 | Loss: 0.8403
Epoch: 037/040 | Train: 71.17% | Validation: 56.78%
Time elapsed: 47.88 min
Epoch: 038/040 | Batch 0000/0175 | Loss: 0.7887
Epoch: 038/040 | Batch 0100/0175 | Loss: 1.2266
Epoch: 038/040 | Train: 71.57% | Validation: 56.60%
Time elapsed: 49.18 min
Epoch: 039/040 | Batch 0000/0175 | Loss: 0.7900
Epoch: 039/040 | Batch 0100/0175 | Loss: 0.9850
Epoch: 039/040 | Train: 71.37% | Validation: 57.90%
Time elapsed: 50.47 min
Epoch: 040/040 | Batch 0000/0175 | Loss: 0.7807
Epoch: 040/040 | Batch 0100/0175 | Loss: 0.8716
Epoch: 040/040 | Train: 69.35% | Validation: 55.66%
Time elapsed: 51.81 min
Total Training Time: 51.81 min
Test accuracy 55.17%
```

```
plot_training_loss(minibatch_loss_list=minibatch_loss_list_rms,
                   num_epochs=NUM_EPOCHS,
                   iter_per_epoch=len(train_loader),
                   results_dir=None,
                   averaging_iterations=20)
```

```
plt.show()
```

```
plot_accuracy(train_acc_list=train_acc_list_rms,
              valid_acc_list=valid_acc_list_rms,
              results_dir=None)
```

```
# plt.ylim([80, 100])
plt.show()
```



```
class CNN2RMS(torch.nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.features = torch.nn.Sequential(
            # Conv 1
            torch.nn.Conv2d(3, 16, kernel_size=3, padding="same"), # output 16 - 3 +
            # , stride=4, padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2), # 16 / 2 => output 8

            # Conv 2
            torch.nn.Conv2d(16, 32, kernel_size=2, padding="same"), # output 7 - 2 +
            # , padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2), #output 8 / 2 => output 4

            torch.nn.Conv2d(32, 64, kernel_size=2, padding="same"), # output 7 - 2 +
            # , padding=2),
            torch.nn.ReLU(inplace=True),
            torch.nn.MaxPool2d(kernel_size=2) #output 8 / 2 => output 2
```

)

```
self.classifier = torch.nn.Sequential(
    torch.nn.Linear(64*2*2, 100),
    torch.nn.ReLU(inplace=True),
    torch.nn.Linear(100, num_classes),
)
```

```
def forward(self, x):
    x = self.features(x)
    x = torch.flatten(x, 1)
    # print(x.size())
    logits = self.classifier(x)
    return logits
```

```
model2_rms = CNN2RMS(num_classes=10)
```

```
model2_rms = model2_rms.to(DEVICE)
```

```
print(summary(model2_rms, (3, 16, 16)))
```



Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 16, 16]	448
ReLU-2	[-1, 16, 16, 16]	0
MaxPool2d-3	[-1, 16, 8, 8]	0
Conv2d-4	[-1, 32, 8, 8]	2,080
ReLU-5	[-1, 32, 8, 8]	0
MaxPool2d-6	[-1, 32, 4, 4]	0
Conv2d-7	[-1, 64, 4, 4]	8,256
ReLU-8	[-1, 64, 4, 4]	0
MaxPool2d-9	[-1, 64, 2, 2]	0
Linear-10	[-1, 100]	25,700
ReLU-11	[-1, 100]	0
Linear-12	[-1, 10]	1,010

```
=====
Total params: 37,494
Trainable params: 37,494
Non-trainable params: 0
=====
```

```
Input size (MB): 0.00
Forward/backward pass size (MB): 0.12
Params size (MB): 0.14
Estimated Total Size (MB): 0.27
=====
```

```
None
```

```
optimizer_rms2 = torch.optim.RMSprop(model2_rms.parameters())
scheduler_rms2 = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer_rms2,
                                                                factor=0.1,
```

```
mode='max',  
verbose=True)
```

```
minibatch_loss_list_rms2, train_acc_list_rms2, valid_acc_list_rms2 = train_model(  
    model=model2_rms,  
    num_epochs=NUM_EPOCHS,  
    train_loader=train_loader,  
    valid_loader=valid_loader,  
    test_loader=test_loader,  
    optimizer=optimizer_rms2,  
    device=DEVICE,  
    scheduler=None,  
    scheduler_on='valid_acc',  
    logging_interval=100)
```







```
Epoch: 036/040 | Batch 0100/0175 | Loss: 1.0481
Epoch: 036/040 | Train: 59.92% | Validation: 54.12%
Time elapsed: 45.01 min
Epoch: 037/040 | Batch 0000/0175 | Loss: 1.0873
Epoch: 037/040 | Batch 0100/0175 | Loss: 1.0833
Epoch: 037/040 | Train: 53.34% | Validation: 48.46%
Time elapsed: 46.26 min
Epoch: 038/040 | Batch 0000/0175 | Loss: 1.3300
Epoch: 038/040 | Batch 0100/0175 | Loss: 1.1263
Epoch: 038/040 | Train: 61.58% | Validation: 54.84%
Time elapsed: 47.51 min
Epoch: 039/040 | Batch 0000/0175 | Loss: 1.1002
Epoch: 039/040 | Batch 0100/0175 | Loss: 1.0842
Epoch: 039/040 | Train: 63.25% | Validation: 56.04%
Time elapsed: 48.76 min
Epoch: 040/040 | Batch 0000/0175 | Loss: 1.0204
Epoch: 040/040 | Batch 0100/0175 | Loss: 1.1030
Epoch: 040/040 | Train: 63.51% | Validation: 55.92%
```

```
plot_training_loss(minibatch_loss_list=minibatch_loss_list_rms2,
                   num_epochs=NUM_EPOCHS,
                   iter_per_epoch=len(train_loader),
                   results_dir=None,
                   averaging_iterations=20)
```

```
plt.show()
```

```
plot_accuracy(train_acc_list=train_acc_list_rms2,
              valid_acc_list=valid_acc_list_rms2,
              results_dir=None)
```

```
# plt.ylim([80, 100])
plt.show()
```



```
import pandas as pd
```

```
results = pd.DataFrame({"Number of Parameters": [54838, 37494], "Accuracy": [70, 63]})
```

```
results
```



	Number of Parameters	Accuracy
<b>CNN1</b>	54838	70
<b>CNN2</b>	37494	63