

- ✓ Pranav Submission
- ✓ 1a - Two parameters need to be learned, they are  $w$  &  $b$

$$1b) y = x_1 + wx_2 + b.$$

$$L = \frac{1}{2n} \sum (\hat{y}^{(i)} - y^{(i)})^2$$

Considering it as scalar for the purpose of derivation.

$$L = \frac{1}{2n} (\hat{y} - y)^2$$

$$L = \frac{1}{2n} (x_1 + wx_2 + b - y)^2$$

$$\frac{\partial L}{\partial w} = \frac{\partial}{\partial w} \cdot \frac{1}{2n} \cdot (x_1 + wx_2 + b - y)^2$$

Applying chain rule:  $\frac{\partial}{\partial x} f(g(x)) = f'(g(x)) \cdot g'(x)$ .

$$\text{Since } \frac{d}{dx} x^2 = 2x$$

$$= \frac{1}{2n} \cdot 2 \cdot (x_1 + wx_2 + b - y) \cdot \frac{\partial}{\partial w} (x_1 + wx_2 + b - y)$$

$$= \frac{1}{n} \cdot (x_1 + wx_2 + b - y) \cdot \frac{\partial}{\partial w} (x_1 + wx_2 + b - y)$$

Since  $\frac{d}{dx} ax = a$  &  $\frac{d}{dx} (\text{constant}) = 0$ .

$$= \frac{1}{n} (x_1 + wx_2 + b - y) \cdot x_2 \quad \left\{ \text{Since } \frac{\partial}{\partial w} (wx_2) = x_2 \right.$$

$$\frac{\partial L}{\partial w} = \frac{1}{n} (\hat{y} - y) \cdot x_2$$



$$1c) y = x_1 + wx_2 + b.$$

$$L = \frac{1}{2n} \sum (\hat{y}(i) - y(i))^2.$$

Considering it as scalar for the purpose of derivation

$$L = \frac{1}{2n} (\hat{y} - y)^2.$$

$$= \frac{1}{2n} (x_1 + wx_2 + b - y)^2.$$

$$\text{App } \frac{\partial L}{\partial b} = \frac{\partial}{\partial b} \frac{1}{2n} (x_1 + wx_2 + b - y)^2.$$

Applying chain rule:  $\frac{\partial}{\partial x} f(g(x)) = f'(g(x)) \cdot g'(x)$ .

$$\text{Since } \frac{d}{dx} (x^2) = 2x.$$

$$= \frac{1}{2n} \cdot 2 \cdot (x_1 + wx_2 + b - y) \cdot \frac{\partial}{\partial b} (x_1 + wx_2 + b - y)$$

$$\text{Since } \frac{d}{dx} (x) = 1 \text{ \& } \frac{d}{dx} (\text{constant}) = 0.$$

$$= \frac{1}{n} \cdot (x_1 + wx_2 + b - y) \cdot 1.$$

$$\frac{\partial L}{\partial b} = \frac{1}{n} (\hat{y} - y).$$

$$\frac{\partial L}{\partial b} = \frac{1}{n} (\hat{y} - y).$$

1d) 1. Initialize  $w := 0 \in \mathbb{R}^m$ ,  $b := 0$ .

2. For every training epoch:

A. For every  $\langle x_1^{(i)}, x_2^{(i)}, y^{(i)} \rangle \in D$ .

a)  $\hat{y}^{(i)} = \sigma(x_1^{(i)} + w x_2^{(i)} + b)$  { Activation function  $\sigma(x) = x$  }

b)  $\frac{\partial L}{\partial w} = \frac{1}{n} (\hat{y}^{(i)} - y^{(i)}) x_2^{(i)}$

$\frac{\partial L}{\partial b} = \frac{1}{n} (\hat{y}^{(i)} - y^{(i)})$

c)  $w := w + \eta \times \left( \frac{1}{n} (\hat{y}^{(i)} - y^{(i)}) x_2^{(i)} \right)$

$b := b + \eta \times \left( \frac{1}{n} (\hat{y}^{(i)} - y^{(i)}) \right)$

learning rate  $\uparrow$  negative gradient

## ✓ 1e solution

```
import pandas as pd
import matplotlib.pyplot as plt
import torch
%matplotlib inline
```

```
df = pd.read_csv('./linreg-data.csv', index_col=0)
```

```
# Assign features and target
```

```
X1 = torch.tensor(df[['x1']].values, dtype=torch.float)
X2 = torch.tensor(df[['x2']].values, dtype=torch.float)
```

```

y = torch.tensor(df['y'].values, dtype=torch.float)

# Shuffling & train/test split

torch.manual_seed(123)
shuffle_idx = torch.randperm(y.size(0), dtype=torch.long)

X1, X2, y = X1[shuffle_idx], X2[shuffle_idx], y[shuffle_idx]

percent70 = int(shuffle_idx.size(0)*0.7)

X1_train, X1_test = X1[shuffle_idx[:percent70]], X1[shuffle_idx[percent70:]]
X2_train, X2_test = X2[shuffle_idx[:percent70]], X2[shuffle_idx[percent70:]]
y_train, y_test = y[shuffle_idx[:percent70]], y[shuffle_idx[percent70:]]

# Normalize (mean zero, unit variance)

mu1, sigma1 = X1_train.mean(dim=0), X1_train.std(dim=0)
mu2, sigma2 = X2_train.mean(dim=0), X2_train.std(dim=0)
X1_train = (X1_train - mu1) / sigma1
X1_test = (X1_test - mu1) / sigma1

X2_train = (X2_train - mu2) / sigma2
X2_test = (X2_test - mu2) / sigma2

X1_train.shape

⇒ torch.Size([700, 1])

class LinearRegression():
    def __init__(self, num_features):
        self.num_features = num_features
        self.weights = torch.zeros(num_features, 1,
                                    dtype=torch.float)
        self.bias = torch.zeros(1, dtype=torch.float)

    def forward(self, x1, x2):
        netinputs = torch.add(torch.add(x1, torch.mm(x2, self.weights)), self.bias)
        activations = netinputs
        return activations.view(-1)

    def backward(self, x1, x2, yhat, y):

        grad_loss_yhat = 2*(yhat - y)

        grad_yhat_weights = x2
        grad_yhat_bias = 1.

        # Chain rule: inner times outer
        grad_loss_weights = torch.mm(grad_yhat_weights.t(),

```

```

grad_loss_yhat.view(-1, 1)) / y.size(0)

grad_loss_bias = torch.sum(grad_yhat_bias*grad_loss_yhat) / y.size(0)

# return negative gradient
return (-1)*grad_loss_weights, (-1)*grad_loss_bias

#####
##### Training and evaluation wrappers
#####

def loss(yhat, y):
    return torch.mean((yhat - y)**2)

def train(model, x1, x2, y, num_epochs, learning_rate=0.01):
    cost = []
    for e in range(num_epochs):

        ##### Compute outputs #####
        yhat = model.forward(x1, x2)

        ##### Compute gradients #####
        negative_grad_w, negative_grad_b = model.backward(x1, x2, yhat, y)

        ##### Update weights #####
        model.weights += learning_rate * negative_grad_w
        model.bias += learning_rate * negative_grad_b

        ##### Logging #####
        yhat = model.forward(x1, x2) # not that this is a bit wasteful here
        curr_loss = loss(yhat, y)
        print('Epoch: %03d' % (e+1), end='')
        print(' | MSE: %.5f' % curr_loss)
        cost.append(curr_loss)

    return cost

model = LinearRegression(num_features=X1_train.size(1))
cost = train(model,
              X1_train, X2_train, y_train,
              num_epochs=100,
              learning_rate=0.05)

```




```
Epoch: 049 | MSE: 372.00595  
Epoch: 050 | MSE: 372.05490  
Epoch: 051 | MSE: 372.04755  
Epoch: 052 | MSE: 372.04160  
Epoch: 053 | MSE: 372.03677  
Epoch: 054 | MSE: 372.03287  
Epoch: 055 | MSE: 372.02972  
Epoch: 056 | MSE: 372.02713  
Epoch: 057 | MSE: 372.02505  
Epoch: 058 | MSE: 372.02341  
Epoch: 059 | MSE: 372.02203  
Epoch: 060 | MSE: 372.02090  
Epoch: 061 | MSE: 372.02005  
Epoch: 062 | MSE: 372.01929  
Epoch: 063 | MSE: 372.01874  
Epoch: 064 | MSE: 372.01825  
Epoch: 065 | MSE: 372.01782  
Epoch: 066 | MSE: 372.01749  
Epoch: 067 | MSE: 372.01730  
Epoch: 068 | MSE: 372.01709  
Epoch: 069 | MSE: 372.01691  
Epoch: 070 | MSE: 372.01678  
Epoch: 071 | MSE: 372.01669  
Epoch: 072 | MSE: 372.01657  
Epoch: 073 | MSE: 372.01651  
Epoch: 074 | MSE: 372.01642  
Epoch: 075 | MSE: 372.01642  
Epoch: 076 | MSE: 372.01633  
Epoch: 077 | MSE: 372.01633  
Epoch: 078 | MSE: 372.01630  
Epoch: 079 | MSE: 372.01624  
Epoch: 080 | MSE: 372.01630  
Epoch: 081 | MSE: 372.01627  
Epoch: 082 | MSE: 372.01624  
Epoch: 083 | MSE: 372.01624  
Epoch: 084 | MSE: 372.01624  
Epoch: 085 | MSE: 372.01620  
Epoch: 086 | MSE: 372.01624  
Epoch: 087 | MSE: 372.01624  
Epoch: 088 | MSE: 372.01624  
Epoch: 089 | MSE: 372.01620  
Epoch: 090 | MSE: 372.01620  
Epoch: 091 | MSE: 372.01624  
Epoch: 092 | MSE: 372.01624  
Epoch: 093 | MSE: 372.01620  
Epoch: 094 | MSE: 372.01620  
Epoch: 095 | MSE: 372.01624  
Epoch: 096 | MSE: 372.01617  
Epoch: 097 | MSE: 372.01620  
Epoch: 098 | MSE: 372.01624  
Epoch: 099 | MSE: 372.01620  
Epoch: 100 | MSE: 372.01620
```

## ✓ Evaluating & Plotting



```
plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()
```



Start coding or [generate](#) with AI.



```
Train MSE: 372.01620
Test MSE: 409.19485
```

```
print('Weights', model.weights)
print('Bias', model.bias)
```



```
Weights tensor([[37.8872]])
Bias tensor([-0.5464])
```

1f The second weight and the bias terms are both similar to the model in the lecture video, however even in the first video the final weight of  $x_1$  was 0.36, the value we have for  $x_1$  is 1, so the accuracies were comparable to the model in lecture video

## ✓ Bonus Question

Learning rate - 0.001


```
model = LinearRegression(num_features=X1_train.size(1))  
cost = train(model,  
              X1_train, X2_train, y_train,  
              num_epochs=100,  
              learning_rate=0.001)
```



```
Epoch: 043 | MSE: 1579.30396  
Epoch: 044 | MSE: 1574.48645  
Epoch: 045 | MSE: 1569.68835  
Epoch: 046 | MSE: 1564.90906  
Epoch: 047 | MSE: 1560.14929  
Epoch: 048 | MSE: 1555.40820  
Epoch: 049 | MSE: 1550.68604  
Epoch: 050 | MSE: 1545.98291
```

```
Epoch: 093 | MSE: 1350.54400
Epoch: 094 | MSE: 1356.60022
Epoch: 095 | MSE: 1352.67151
Epoch: 096 | MSE: 1348.75842
Epoch: 097 | MSE: 1344.86084
Epoch: 098 | MSE: 1340.97900
Epoch: 099 | MSE: 1337.11255
Epoch: 100 | MSE: 1333.26135
```

```
plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()
```




Learning rate - 0.01

```
model = LinearRegression(num_features=X1_train.size(1))
cost = train(model,
              X1_train, X2_train, y_train,
              num_epochs=100,
              learning_rate=0.001)
```

```
Epoch: 001 | MSE: 1800.05798
Epoch: 002 | MSE: 1794.35962
Epoch: 003 | MSE: 1788.68408
Epoch: 004 | MSE: 1783.03125
Epoch: 005 | MSE: 1777.40088
Epoch: 006 | MSE: 1771.79285
Epoch: 007 | MSE: 1766.20752
Epoch: 008 | MSE: 1760.64429
Epoch: 009 | MSE: 1755.10327
Epoch: 010 | MSE: 1749.58423
Epoch: 011 | MSE: 1744.08728
Epoch: 012 | MSE: 1738.61230
Epoch: 013 | MSE: 1733.15930
Epoch: 014 | MSE: 1727.72791
Epoch: 015 | MSE: 1722.31824
```

Epoch: 016	MSE: 1716.93005
Epoch: 017	MSE: 1711.56360
Epoch: 018	MSE: 1706.21826
Epoch: 019	MSE: 1700.89441
Epoch: 020	MSE: 1695.59180
Epoch: 021	MSE: 1690.31030
Epoch: 022	MSE: 1685.05005
Epoch: 023	MSE: 1679.81055
Epoch: 024	MSE: 1674.59216
Epoch: 025	MSE: 1669.39429
Epoch: 026	MSE: 1664.21729
Epoch: 027	MSE: 1659.06104
Epoch: 028	MSE: 1653.92542
Epoch: 029	MSE: 1648.81018
Epoch: 030	MSE: 1643.71558
Epoch: 031	MSE: 1638.64111
Epoch: 032	MSE: 1633.58679
Epoch: 033	MSE: 1628.55273
Epoch: 034	MSE: 1623.53870
Epoch: 035	MSE: 1618.54480
Epoch: 036	MSE: 1613.57068
Epoch: 037	MSE: 1608.61646
Epoch: 038	MSE: 1603.68213
Epoch: 039	MSE: 1598.76746
Epoch: 040	MSE: 1593.87256
Epoch: 041	MSE: 1588.99683
Epoch: 042	MSE: 1584.14075
Epoch: 043	MSE: 1579.30396
Epoch: 044	MSE: 1574.48645
Epoch: 045	MSE: 1569.68835
Epoch: 046	MSE: 1564.90906
Epoch: 047	MSE: 1560.14929
Epoch: 048	MSE: 1555.40820
Epoch: 049	MSE: 1550.68604
Epoch: 050	MSE: 1545.98291
Epoch: 051	MSE: 1541.29834
Epoch: 052	MSE: 1536.63245
Epoch: 053	MSE: 1531.98535
Epoch: 054	MSE: 1527.35681
Epoch: 055	MSE: 1522.74683
Epoch: 056	MSE: 1518.15503
Epoch: 057	MSE: 1513.58167
Epoch: 058	MSE: 1509.02625

```
plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()
```



Learning rate - 0.1

```
model = LinearRegression(num_features=X1_train.size(1))
cost = train(model,
              X1_train, X2_train, y_train,
              num_epochs=100,
              learning_rate=0.1)
```



```
Epoch: 001 | MSE: 1290.28015
Epoch: 002 | MSE: 960.12476
Epoch: 003 | MSE: 748.67462
Epoch: 004 | MSE: 613.24976
Epoch: 005 | MSE: 526.51587
Epoch: 006 | MSE: 470.96671
Epoch: 007 | MSE: 435.38974
Epoch: 008 | MSE: 412.60419
Epoch: 009 | MSE: 398.01108
Epoch: 010 | MSE: 388.66476
Epoch: 011 | MSE: 382.67889
Epoch: 012 | MSE: 378.84521
Epoch: 013 | MSE: 376.38986
Epoch: 014 | MSE: 374.81735
Epoch: 015 | MSE: 373.81021
Epoch: 016 | MSE: 373.16519
Epoch: 017 | MSE: 372.75204
Epoch: 018 | MSE: 372.48746
Epoch: 019 | MSE: 372.31802
Epoch: 020 | MSE: 372.20950
Epoch: 021 | MSE: 372.14001
Epoch: 022 | MSE: 372.09549
Epoch: 023 | MSE: 372.06702
Epoch: 024 | MSE: 372.04871
Epoch: 025 | MSE: 372.03702
Epoch: 026 | MSE: 372.02954
Epoch: 027 | MSE: 372.02478
Epoch: 028 | MSE: 372.02170
Epoch: 029 | MSE: 372.01971
```

```
Epoch: 030 | MSE: 372.01840
Epoch: 031 | MSE: 372.01764
Epoch: 032 | MSE: 372.01715
Epoch: 033 | MSE: 372.01678
Epoch: 034 | MSE: 372.01657
Epoch: 035 | MSE: 372.01642
Epoch: 036 | MSE: 372.01639
Epoch: 037 | MSE: 372.01630
Epoch: 038 | MSE: 372.01630
Epoch: 039 | MSE: 372.01624
Epoch: 040 | MSE: 372.01624
Epoch: 041 | MSE: 372.01624
Epoch: 042 | MSE: 372.01620
Epoch: 043 | MSE: 372.01617
Epoch: 044 | MSE: 372.01624
Epoch: 045 | MSE: 372.01617
Epoch: 046 | MSE: 372.01620
Epoch: 047 | MSE: 372.01620
Epoch: 048 | MSE: 372.01620
Epoch: 049 | MSE: 372.01620
Epoch: 050 | MSE: 372.01620
Epoch: 051 | MSE: 372.01624
Epoch: 052 | MSE: 372.01620
Epoch: 053 | MSE: 372.01620
Epoch: 054 | MSE: 372.01620
Epoch: 055 | MSE: 372.01624
Epoch: 056 | MSE: 372.01620
Epoch: 057 | MSE: 372.01624
Epoch: 058 | MSE: 372.01617
```

```
plt.plot(range(len(cost)), cost)
plt.ylabel('Mean Squared Error')
plt.xlabel('Epoch')
plt.show()
```

